

OISA: A Formalism-Agnostic Orchestration Architecture for Composing Published Immune Models Without Modification

Nathan Foulquier

LBAI — Laboratoire de Biologie des Lymphocytes Auto-Immuns, Inserm U1227, Université de Bretagne Occidentale
Centre de Données Cliniques, CHU de Brest

Brest, France

Email: nathan.foulquier@univ-brest.fr

Abstract—Computational immunology has produced high-quality mechanistic models of individual immune compartments, yet composing ODE and agent-based models (ABMs) across formalism boundaries still requires bespoke re-engineering. Existing standards (SBML, CellML) address intra-formalism portability; Vivarium introduced formalism-agnostic composition but provides no standardised runtime uncertainty quantification, model-derived transfer lags, or biological plausibility constraint engine. We propose OISA — the Orchestrated Immune Simulation Architecture — comprising (i) the ISSL, a formalism-agnostic JSON-LD checkpoint format with embedded runtime UQ (uncertainty bounds propagated at every tick via the `ci_95` schema field); (ii) a declarative configuration graph with support for model-derived transfer lags on edges; and (iii) an orchestrator engine with global clock synchronisation, causal DAG resolution, and biological plausibility enforcement.

We validate OISA by coupling two independently published influenza models — Miao *et al.* 2010 ODE (SBML BIOMD0000000546, unmodified) and the full spatial Sego *et al.* 2020 ABM (CompuCell3D 4.8.0, 12 steppables unmodified on a $90 \times 90 \times 2$ grid) — with zero changes to either model’s source code. Adapter layers total ~300 lines. The ABM runs as a rolling ensemble of $N = 5$ parallel CC3D instances (concurrent subprocesses, distinct MersenneTwister seeds); runtime uncertainty bounds — reported as *ensemble range* across $N = 5$ *stochastic CC3D replicates* (sample min–max, see §V-A) — originate in the stochastic ABM and propagate through the ODE at each checkpoint. Over a 14-day simulation (56 checkpoints), the coupled system reproduces viral kinetics consistent with Miao 2010 (peak $V = 9.0 \times 10^6$ copies/mL at day 2.25, clearance to $< 0.1\%$ of peak by day 13.75) and spatial immune recruitment (median $n_{\text{immune}} = 6$ [ens.range: 4–7] at day 1, growing to 50 [ens.range: 40–53] at day 13), with immune temporal lag and CTL-mediated clearance acceleration emerging from the bidirectional coupling. This validation targets orchestration correctness — signal routing, causal ordering, and runtime UQ propagation — not biological prediction of the coupled system, which has not been validated against simultaneous in-vivo measurements.

Index Terms—immune digital twins; multi-scale modelling; model composition; agent-based models; ordinary differential equations; interoperability; computational immunology; CompuCell3D.

I. INTRODUCTION

Immune-mediated diseases are inherently multi-compartmental phenomena. Influenza infection unfolds across scales: within a single tissue, viral replication dynamics interact with innate immune cytokine signalling and adaptive CTL recruitment, each operating on different timescales and representable in different formalisms. ODEs are natural for systemic viral kinetics and antibody dynamics, while ABMs are necessary where stochastic single-cell fate decisions dominate, as in tissue-level immune recruitment. The literature has consequently produced a collection of high-quality compartment-specific models that cannot communicate with one another.

Three bodies of prior work address adjacent problems. The COMBINE standards ecosystem (SBML [1], CellML [2], SED-ML [3]) provides portable representations for individual models but addresses intra-formalism composition only: models must be expressed in the same declarative format, ruling out ABMs written in Python or Julia. Vivarium [4] introduced formalism-agnostic composition with hierarchical stores and dynamic topology (discussed in §II-A), but provides no standardised runtime UQ propagation, model-derived transfer lags, or biological plausibility constraint engine. The CURE guidelines [5] define what properties a credible model should exhibit — Credibility, Understandability, Reproducibility, Extensibility — but prescribe no runtime infrastructure for heterogeneous composition (note: [5] is currently available as a preprint pending peer review). The cross-formalism, multi-timescale coordination problem remains open.

We propose OISA — the Orchestrated Immune Simulation Architecture — to fill this gap. OISA provides three components: the ISSL formalism-agnostic state interface (§IV-A), a declarative configuration graph (§IV-B), and an orchestrator engine (§IV-C). We demonstrate the architecture by coupling two independently published influenza models, using the full spatial Sego 2020 CompuCell3D ABM without any modification to its 12 biological steppables, and evaluate causal

ordering, signal flow correctness, and biological plausibility in §V.

The contributions of this work are:

- 1) **The ISSL**, a formalism-agnostic checkpoint format with embedded provenance and runtime uncertainty quantification (ensemble bounds propagated at each tick from stochastic ABM through deterministic ODE via the `ci_95` schema field);
- 2) **A declarative configuration graph** with support for model-derived transfer lags on edges, where a lightweight transfer model computes the signal delay dynamically rather than using a fixed constant;
- 3) **A runtime orchestrator** with global clock synchronisation, causal DAG resolution, and biological plausibility constraint enforcement;
- 4) **Empirical demonstration** that two independently published models (Miao 2010 SBML ODE + Sego 2020 CC3D spatial ABM) can be composed with zero source modifications using ~300 lines of adapter code, with runtime UQ propagated end-to-end.

II. RELATED WORK

A. Multi-Formalism Immune Simulation Frameworks

Several frameworks have addressed multi-scale composition in computational biology. Vivarium [4] introduced a port-based, formalism-agnostic composition interface with hierarchical stores, dynamic topology changes at runtime, and a mature Python ecosystem. OISA shares Vivarium’s design principle of formalism-agnostic composition but addresses three capabilities not present in the Vivarium architecture: (i) a standardised inter-model signal format (ISSL) embedding runtime uncertainty quantification (ensemble bounds propagated at each checkpoint from a rolling ABM ensemble through the ODE via the `ci_95` schema field), (ii) model-derived transfer lags on edges (where a lightweight transfer model computes the signal delay dynamically rather than using a fixed constant), and (iii) a biological plausibility constraint engine operating at the composition level (mass conservation, parameter bounds, and divergence monitoring across coupled models). Vivarium’s store-based data sharing and dynamic topology management are complementary capabilities not replicated by OISA.

PhysiCell [6] and PhysiBoSS [7] extend ABM with Boolean signalling layers but operate within a single formalism. The rapid community-driven SARS-CoV-2 tissue simulator [8] — from which Sego *et al.* 2020 [9] derives — demonstrated that tissue-scale ABMs can be developed and shared rapidly, but required a uniform CompuCell3D substrate for all constituent models; coupling to external ODE models required bespoke adapter code outside the framework. Miao *et al.* 2010 [10] provided a rigorously calibrated ODE model of murine influenza kinetics embedded in the BioModels database (BIOMD0000000546); no standard mechanism exists to couple it to published tissue ABMs without rewriting either model. OISA demonstrates that such coupling is achievable by

adding only a thin Emit/Accept interface layer to each model, with the full CC3D spatial ABM running unmodified as an independent process.

B. Simulation Interoperability Standards

The COMBINE standards ecosystem [11] has achieved significant progress on individual model portability. SBML [1] has accumulated over 1,000 curated models in BioModels; CellML [2] has analogous achievements in cardiac and physiological modelling; SED-ML [3] provides a standard for encoding simulation experiments. Table I compares these standards against OISA. The critical distinction is architectural: SBML, CellML, and NeuroML address *intra-formalism* interoperability — exchange between tools supporting the same format — while OISA addresses *inter-formalism* interoperability: composition of models using fundamentally different formalisms, timescales, and output semantics. SBML’s hierarchical composition package (comp) extends composition within SBML but cannot incorporate an ABM without first rewriting it in SBML, which is both impractical and biologically inappropriate for stochastic cellular processes.

C. Immune Digital Twins

The concept of the immune digital twin (IDT) was formalised by Laubenbacher *et al.* [12] as a continuously-recalibrated patient-specific computational model of the immune system. Subsequent reviews [13], [14] have identified interoperability as the primary technical barrier to clinical deployment, a conclusion echoed by the National Academies [15], by model integration consortia [16], and by multi-scale digital twin frameworks in adjacent domains [17]. OISA is the first runtime orchestration architecture designed specifically for heterogeneous immune model composition, providing the coordination infrastructure that these reviews identify as lacking.

III. PROBLEM FORMALISATION

A. Formal Model Definition

We define a composable simulation model M as a triple $M = (S, \text{Emit}, \text{Accept})$, where S is the model’s internal state space, $\text{Emit} : S \times T \rightarrow \text{ISSL}$ maps the current state and simulation time to a structured ISSL record, and $\text{Accept} : \text{ISSL} \rightarrow S$ updates the model state in response to an incoming inter-model signal. Any model — ODE, ABM, hybrid, or neural surrogate — that implements Emit and Accept is OISA-compliant without any other modification. Formalism-agnosticism follows directly: since the ISSL is the only communication channel, models need not know each other’s internal representation, only the schema of the signal they receive. This design principle — letting each biological process be modelled in its natural formalism and solving the composition problem at the interface — avoids forcing biologically inappropriate representations on individual compartments.

TABLE I

COMPARISON OF SBML, CELLML, NEUROML/LEMS, AND OISA ACROSS CAPABILITIES RELEVANT TO MULTI-ORGAN IMMUNE SIMULATION. “—” = OUT OF ARCHITECTURAL SCOPE (OISA IS A RUNTIME ORCHESTRATION ARCHITECTURE, NOT A PORTABLE MODEL REPRESENTATION FORMAT). RUNTIME UQ: ENSEMBLE BOUNDS (REPORTED IN THE `ci_95` SCHEMA FIELD) ORIGINATE IN THE STOCHASTIC ABM ENSEMBLE ($N=5$ PARALLEL CC3D INSTANCES; WITH $N=5$, REPORTED BOUNDS CORRESPOND TO SAMPLE MIN-MAX, NOT STABLE PERCENTILE ESTIMATES) AND PROPAGATE THROUGH THE DETERMINISTIC ODE VIA TRIPLE INTEGRATION AT EACH CHECKPOINT (§V-A). “SBML L3” REFERS TO THE FORMAT SPECIFICATION ALONE; SED-ML [3] PROVIDES CO-SIMULATION CAPABILITIES INCLUDING HETEROGENEOUS TIME-STEPPING FOR SBML-COMPATIBLE MODELS BUT DOES NOT EXTEND TO ABM COMPOSITION.

Capability	SBML L3	CellML 2.0	NeuroML/LEMS	OISA (ISSL)
Portable model representation	✓	✓	✓	—
ODE / biochemical networks	✓	✓	neurons only	via <code>Emit()</code>
Agent-based models	×	×	×	via <code>Emit()</code>
Heterogeneous Δt between models	×	×	×	✓
Model-derived transfer lag on edges	×	×	×	✓
Inter-formalism composition (ODE + ABM)	×	×	×	✓
Runtime UQ propagation across models	×	×	×	✓
OBO ontology annotation	✓	✓	✓	✓ + wildcard
ABM scaling factor declaration	×	×	×	✓
Requires model rewrite to adopt	✓	✓	✓	×

B. The Composition Problem

A multi-model composition is a directed graph $G=(V, E)$ where each vertex v is a model M_v and each directed edge (u, v) specifies that a subset of u ’s `export_signals` are routed to v ’s `Accept` function, with optional transfer-model lags. The composition problem reduces to three coordination challenges:

- 1) **Heterogeneous time steps.** A global simulation clock (GSimT) must coordinate execution such that no model receives a signal purporting to come from a future GSimT tick. In the influenza reference implementation, the viral dynamics ODE steps at 6 h intervals while the tissue immune ABM steps at 24 h intervals (72 Monte Carlo Steps at 20 min/MCS); the GSimT tick is set to the GCD of all model Δt_i values (here 6 h).
- 2) **Stochastic–deterministic reconciliation.** Signals from a stochastic ABM to a deterministic ODE must carry uncertainty information. The orchestrator normalises all inter-model signals to (median, uncertainty bounds, unit) before routing, stored in the `ci_95` schema field: for stochastic models, bounds are computed from a rolling ensemble of N parallel instances (§V-A; with $N=5$, these correspond to sample min–max rather than stable percentile estimates); for deterministic models, bounds are derived by propagating the input uncertainty through the model equations.
- 3) **Causal ordering with feedback.** The immune model emits cytokine signals that drive viral clearance; viral load in turn drives immune recruitment — a feedback loop. The orchestrator resolves this by applying a one-tick delay on back-edges of the DAG, preventing circular dependency while preserving biological causality at the GSimT timescale.

IV. THE OISA ARCHITECTURE

A. The Internal Simulation State Log (ISSL)

The ISSL is a JSON-LD document emitted by a model at each checkpoint, structured in six sections (Table II). JSON-

TABLE II

ISSL SECTION INVENTORY. THE `EXPORT_SIGNALS` FIELD ALWAYS CARRIES THE BIOLOGICALLY SCALED VALUE, ENSURING DOWNSTREAM MODELS RECEIVE CORRECT QUANTITIES REGARDLESS OF SOURCE FORMALISM.

ISSL section	Content / parsability note
envelope	Model ID, version, GSimT timestamp, schema URI, formalism tag. Orchestrator validates <code>schema_uri</code> ; reads formalism before processing outputs.
continuous_state	Running entity populations: count (scaled), unit, <code>ci_95</code> (ensemble bounds). label is model-local; orchestrator maps via <code>signal_id</code> routing.
discrete_events	Punctual events: cell death, cytokine peaks. <code>event_type</code> from controlled vocabulary; ABM events include <code>n_realisations</code> .
export_signals	Inter-compartmental signals available for routing. <code>signal_id</code> is the routing key; value in declared unit.
internal_parameters	Kinetic parameter values; provenance is DOI or database URI linking to the calibration source.
watchdog	Model health: status, <code>divergence_score</code> . Values $>0.15 \Rightarrow$ PAUSE.

LD enables both human readability and machine-parsable semantic annotation via linked data context, allowing orchestrator components to parse ISSL records without prior knowledge of the model’s internal variable naming conventions.

Box 1 — ISSL excerpt, Miao 2010 ODE adapter (day 1 checkpoint, GSimT = 86,400 s):

```
{
  "envelope": {
    "model_id": "miao2010_ode",
    "model_version": "BIOMD0000000546",
    "sim_time_s": 86400,
    "formalism": "ODE",
    "schema_uri": "schemas/issl_v1.schema.json"
  },
  "continuous_state": [
    { "label": "Ep", "count": 532048.1, "unit": "cells",
      "ci_95": [530102.3, 533994.0] },
    { "label": "Eps", "count": 39204.6, "unit": "cells",
      "ci_95": [37801.2, 40608.0] },
    { "label": "V", "count": 3.29e5, "unit": "copies/mL",
```

```

"ci_95": [3.11e5, 3.47e5]
},
"export_signals": [
  {
    "signal_id": "miao2010.viral_load",
    "value": 3.29e5, "unit": "copies/mL",
    "ci_95": [3.11e5, 3.47e5], "transfer_lag_s": null
  }
],
"watchdog": {
  "status": "OK", "divergence_score": 0.0
}

```

Note on `sim_time_s`: This field records the clock value *after* the step completes — i.e., the simulation time at which the emitted state is valid.

Box 2 — ISSL excerpt, Sego 2020 CC3D ABM adapter (day 1 checkpoint). Values are from a single illustrative CC3D instance and may differ from the ensemble median ($N = 5$) reported in §V-B4.

```

{
  "envelope": {
    "model_id": "sego2020_abm_cc3d",
    "model_version": "github:covid-tissue-models@5b7e42c+OISA_bridge",
    "sim_time_s": 86400,
    "formalism": "ABM",
    "agent_count": 7,
    "grid": "90x90x2"
  },
  "continuous_state": [
    {
      "label": "n_immune", "count": 7, "unit": "cells",
      "ci_95": [5, 8]
    }
  ],
  "export_signals": [
    {
      "signal_id": "sego2020.immune_cell_count",
      "value": 7, "unit": "cells", "ci_95": [5, 8],
      "transfer_lag_s": null
    },
    {
      "signal_id": "sego2020.recruitment_cytokine",
      "value": 700.0, "unit": "AU", "ci_95": [500.0, 800.0],
      "transfer_lag_s": null
    }
  ],
  "watchdog": {
    "status": "OK", "divergence_score": 0.0
  }
}

```

The SBML file `BIOMD0000000546_model1.xml` is loaded verbatim by `libroadrunner`; no XML modifications are made. The adapter calls `rr.simulate(t0, t1, 2)` and reads species concentrations via the standard `roadrunner` API. For the Sego 2020 CC3D ABM, the adapter launches `CompuCell3D 4.8.0` as an independent subprocess registering all 12 original steppables without modification; the only addition is `OISABridgeSteppable`, which executes every 72 MCS (= 24 h) and handles IPC exclusively — no biological equations are touched.

B. The Configuration Graph

The configuration graph is a YAML or JSON-LD file that fully specifies the composition (Table III). It is the orchestrator's sole input at initialisation; no model needs to know about any other model in the composition.

Influenza coupling graph (excerpt):

```

models:
  - id: miao2010_ode
    formalism: ODE
    executable: models/ode_miao2010/miao2010_adapter.py
    delta_t_s: 21600 # 6 h
  - id: sego2020_abm
    formalism: ABM
    executable: models/abm_sego2020/sego2020_adapter.py
    delta_t_s: 86400 # 24 h (72 MCS x 20 min/MCS)

```





TABLE IV

ORCHESTRATOR COMPONENT INVENTORY. STATUS COLUMN: ✓ IMPLEMENTED AND COVERED BY THE PUBLISHED TEST SUITE (67 TESTS); ○ INTERFACE SPECIFIED, FULL IMPLEMENTATION DEFERRED. COMPONENTS MARKED ✓/○[†] HAVE CORE FUNCTIONALITY IMPLEMENTED AND TESTED (ISSL PARSING, PARAMETER BOUNDS CHECKING, CHECKPOINT SNAPSHOTS, DIVERGENCE THRESHOLD) WHILE SPECIFIC SUB-CAPABILITIES — OOD DETECTION, ROLLBACK RECOVERY, PROV-O URIS, CONSERVATION VIOLATION AUTO-ROLLBACK — ARE SPECIFIED BUT NOT YET IN THE PUBLISHED TEST SUITE.

Component	Responsibility / note	Status
1. ISSL ingestion	JSON-LD parser; SI unit normaliser; OOD detector (Mahalanobis)	✓/○ [†]
2. Temporal scheduler	Priority queue by <code>next_step_due</code> ; heterogeneous Δt	✓
3. Causal resolver	Topological sort; cycle detection; feedback edges get one-tick delay	✓
4. Constraint engine	Mass conservation; parameter bounds; raises <code>CONSTRAINT_VIOLATION</code>	✓/○ [†]
5. State registry	Immutable versioned snapshots per GSimT checkpoint; provenance graph	✓/○ [†]
6. Transfer dispatcher	Invokes transfer model on-demand; maintains <code>SignalQueue</code>	✓
7. Calibration bridge	EHR ingest; triggers recalibration; broadcasts updated priors	○
8. Output aggregator	Merges GIS into checkpoint log; uncertainty propagation	✓
9. Watchdog monitor	Poll model health; pause / resume / rollback	✓/○ [†]

C. The Orchestrator Engine

The orchestrator is a nine-component server process (Table IV). It reads the configuration graph at startup, establishes connections to all model processes, and manages the simulation run.

Temporal scheduling. At each GSimT tick (the GCD of all model Δt_i , here 6h), the scheduler dispatches step commands to models whose `next_step_due` equals the current GSimT. The Sego 2020 CC3D ABM, with $\Delta t = 24\text{h}$ (72 MCS $\times$ 20 min/MCS), receives step commands every fourth tick. The ODE receives a step command at every tick.

Causal resolution. The causal resolver constructs a DAG from the configuration graph and performs topological sorting to determine execution order within each GSimT tick. At each 24h boundary (every fourth tick), the execution order is: (1) `ABM.emit()` — export current immune cell count from CC3D cell inventory; (2) `ODE.accept()` — inject `n_immune` $\rightarrow T_{E_T}$ in SBML; (3) `ODE.step( $\Delta t$ )`; (4) `ODE.emit()` — export viral load. The ODE’s viral load from the previous tick is queued for the ABM’s next `Accept` call (written to `ode_signal.json` before the next CC3D bridge step), implementing the one-tick delay that prevents circular dependency while preserving biological causality.

Calibration bridge (component 7). The calibration bridge has its interface fully specified in the configuration graph schema (Table III, `calibration` element), but its implementation in the reference orchestrator is limited to schema validation of the `calibration` configuration block; full

EHR ingestion and dynamic recalibration are deferred to clinical deployment.

V. VALIDATION

The influenza A viral dynamics coupling (Miao 2010 ODE + Sego 2020 full CC3D ABM) is used as a reference implementation because both models are independently published, their parameters are rigorously calibrated against published experimental data, and their coupling tests a biologically significant feedback loop: viral load drives immune recruitment; CTL cells modulate viral clearance. **No biological novelty is claimed; the coupling serves to validate that OISA correctly orchestrates two heterogeneous published models with zero modification to either.**

A. Validation Setup

Models. The ODE component is BIOMD0000000546 [10], a three-species influenza tissue model (Ep: uninfected epithelial cells; Eps: infected epithelial cells; V: virus) with kinetics calibrated against Murphy *et al.* 1973 murine infection data. The SBML file is downloaded from BioModels and loaded verbatim by `libroadrunner`.

The ABM component is the complete spatial tissue ABM of Sego *et al.* 2020 [9], running in `CompuCell3D` 4.8.0. The simulation occupies a $90 \times 90 \times 2$ voxel grid ($\approx 8,100$ epithelial cell sites; voxel length $4\mu\text{m}$), with six cell types: Medium, Uninfected, Infected, VirusReleasing, Dying, and ImmuneCell (CC3D `typeId=5`). All 12 published steppables are registered without modification: `CellsInitializer`, `ViralReplication`, `ViralInternalization`, `ViralSecretion`, `ImmuneCellKilling`, `Chemotaxis`, `ImmuneCellSeeding`, `SimData`, `CytokineProductionAbsorption`, `ImmuneRecruitment`, `oxidationAgentModel`, `VirusFieldInitializer` (all with the Steppable suffix in the source). The simulation runs at 20 min/MCS; 72 MCS correspond to one 24h GSimT tick; 1,010 MCS correspond to a 14-day run. Three diffusion fields (Virus, cytokine, oxidator) are solved by `DiffusionSolverFE` at each MCS.

The single addition to the CC3D setup is `OISABridgeSteppable` (frequency = 72 MCS), which runs inside the CC3D process but contains no biological equations: it reads the CC3D cell inventory to count ImmuneCell agents (type 5), writes the count to `abm_out.json`, raises an `abm_ready` file sentinel, and blocks until the sentinel is deleted. The adapter writes incoming ODE signals to `ode_signal.json` (non-blocking); the bridge reads and applies these at the next execution. The CC3D process runs as a subprocess of the adapter; no shared memory is used, ensuring process isolation between the ODE solver and the CC3D engine.

Adapter metrics. The total modification to either published model is zero lines (Table V). The Miao 2010 adapter wraps `libroadrunner` (~ 60 lines). The Sego 2020 CC3D adapter manages the subprocess lifecycle and IPC protocol (~ 180 lines),

TABLE V

LINES OF CODE MODIFIED VS. ADDED PER MODEL. THE “ZERO LINES MODIFIED” PROPERTY DIRECTLY DEMONSTRATES OISA’S NON-INVASIVE INTEROPERABILITY CLAIM.

Model	Format	Source	Mod.	Adapter
Miao 2010 ODE	SBML XML	BIOMD0000000546	0	~60
Sego 2020 CC3D ABM	CC3D (Py/XML)	GitHub @5b7e42c	0	~240

complemented by the `OISABridgeSteppable` internal to CC3D (~60 lines); combined ~240 adapter lines with zero Sego 2020 source modifications.

Runtime UQ. The ABM component runs as a rolling ensemble of $N = 5$ parallel CC3D instances (distinct MersenneTwister seeds, isolated IPC directories). The N instances are launched by `Sego2020Ensemble` as concurrent OS subprocesses on separate CPU cores (one `python -m cc3d.run_script` subprocess per seed, scheduled by the OS); the orchestrator blocks on the slowest instance at each 24 h barrier via the IPC sentinel file. Wall-clock cost per tick is therefore $\max(T_i)$ rather than $\sum T_i$, with small overhead from the file-based IPC barrier; the $N \times$ memory cost remains. At each 24 h GSimT boundary, the ensemble adapter aggregates `n_immune` across instances and computes the median and empirical bounds. *Statistical note:* with $N = 5$ instances, the `ci_95` field reports the sample minimum and maximum (empirical p_0 and p_{100}), not stable estimates of the 2.5th and 97.5th percentiles, which would require $N \geq 20$. We retain the schema field name `ci_95` for format consistency but label all reported bounds as “ensemble range [$N=5$]” throughout. The ODE adapter receives the median `n_immune` for its central trajectory and performs three roadrunner integrations per tick (at median, lower-bound, and upper-bound input values), reporting the resulting viral load bounds in its ISSL record. This constitutes runtime UQ propagation: uncertainty originates in the stochastic ABM and is carried through the deterministic ODE at each checkpoint.

Model-derived transfer lags. ISSL export signals carry an optional `transfer_lag_s` field. When non-null, the orchestrator’s `SignalQueue` defers signal injection until the specified delay elapses. A reference implementation (`BloodTransitAdapter`) demonstrates this capability: a first-order ODE for blood transit (Donskoy & Goldschneider 1992) computes $\tau = 1/(\lambda_{\text{seed}} + \lambda_{\text{clear}}) = 4.0$ days as a model-derived lag. The `BloodTransitAdapter` was validated independently (8 dedicated tests) but was not active during the 14-day coupled influenza simulation of §V-B4; the influenza use case uses constant-lag edges.

Signals. Two ISSL signals are routed: `miao2010.viral_load` (copies/mL) from ODE $\rightarrow$ ABM, and `sego2020.immune_cell_count` (`n_immune`, CC3D Immuncell agent count) from ABM $\rightarrow$ ODE.

The ODE $\rightarrow$ ABM coupling operates at the *immune recruitment interface*: the Miao 2010 ODE governs systemic viral kinetics; the Sego 2020 CC3D ABM governs spatial immune cell dynamics. The ODE viral

load is mapped to the CC3D immune recruitment signal (`ImmuneRecruitmentSteppable.totalCytokine`) rather than to the CC3D Virus diffusion field, since viral spread within tissue is already handled by the Sego 2020 steppables at their native spatial resolution.

Derivation of coupling constant
 $\kappa = 3.5 \times 10^{-7} \text{ AU} \cdot \text{mL}/\text{copies}.$ The `ImmuneRecruitmentSteppable` recruits immune cells in proportion to `totalCytokine`, which accumulates from infected-cell secretion at rate $\text{max_ck_secrete_infect} = 10 \times \text{max_ck_secrete_im} = 3.5 \times 10^{-3} \text{ pM/s}$ (`ViralInfectionVTMMModelInputs.py`, line 129). In the Miao 2010 ODE, the infected fraction at quasi-static equilibrium satisfies $\text{Eps}/V \approx \beta_a/\delta_{Es} \approx 1.7 \times 10^{-6} \text{ mL}/\text{copies}$. The per-step cytokine increment contributed by the ODE viral load V is therefore $\Delta \text{totalCytokine} \approx V \cdot 1.7 \times 10^{-6} \cdot 3.5 \times 10^{-3} \cdot \Delta t$. Evaluated at $\Delta t = 86,400 \text{ s}$: $\Delta \text{totalCytokine}/V \approx 5 \times 10^{-4} \text{ pM} \cdot \text{mL}/\text{copies}$. The constant 3.5×10^{-7} (applied per 6 h GSimT tick; four ticks per ABM day) recovers an equivalent daily accumulation $\approx 1.4 \times 10^{-6}$, approximately 2.5 orders of magnitude below the analytical estimate; κ was therefore set empirically to maintain `totalCytokine` within Sego 2020’s functional range at typical peak viral loads (10^6 – 10^7 copies/mL). The discrepancy is absorbed by the unit difference between the ODE’s pM cytokine scale and the ABM’s dimensionless AU accumulator. Preliminary exploration suggests that varying κ by ± 1 order of magnitude shifts the immune onset day by ± 1 –2 days without qualitatively altering the `n_immune` trajectory shape; a formal sensitivity study is reported in §V-B5.

The `sego2020.immune_cell_count` signal sets $T_{ET} = n_{\text{immune}} \times 100 \text{ CTL/mL}$ in the SBML, modulating the CTL killing term $k_E \cdot \text{Eps} \cdot T_{ET}$ (Miao 2010, Table 1; $k_E = 2 \times 10^{-5} \text{ mL} \cdot \text{cell}^{-1} \cdot \text{day}^{-1}$). The scaling factor of 100 CTL/mL per CC3D agent is consistent with published peak CTL densities of 10^3 – 10^6 CTL/mL given a tissue compartment volume of $\approx 0.01 \text{ mL}$. Note that T_{ET} represents adaptive cytotoxic T lymphocytes, while the CC3D Immuncell agents model innate immune effectors (NK-like, cytokine-recruited). Mapping innate spatial agents to an adaptive killing term is a model-level approximation appropriate for demonstrating OISA coupling but not for making biological claims about CTL kinetics.

GSimT configuration. $\text{GCD}(6 \text{ h}, 24 \text{ h}) = 6 \text{ h} \Rightarrow 56$ checkpoints over 14 days. ABM steps at 24 h boundaries; ODE steps at every 6 h tick. Causal order at each 24 h boundary: `ABM.emit()` $\rightarrow$ `ODE.accept()` $\rightarrow$ `ODE.step()` $\rightarrow$ `ODE.emit()` (queued for next ABM IPC cycle).

Test suite. Validation is structured as 67 automated tests (pytest): 23 ODE adapter tests, 20 ABM adapter tests, 8 blood transit transfer model tests, and 16 integration tests. All tests are traceable to published figures or parameter tables; UQ tests verify that `ci_95` fields are non-null when ensemble input is present and that tighter input bounds produce narrower output

TABLE VI

BIOLOGICAL PLAUSIBILITY CHECKS AGAINST PUBLISHED MODEL VALUES. “OISA OUTPUT” IS READ FROM ISSL CHECKPOINTS BY THE ORCHESTRATOR; “PUBLISHED RANGE” IS THE TEST’S ENFORCEMENT TARGET.

Quantity	OISA output	Published range	OK
$V(0)$ viral inoculum	1,000 copies/mL	1,000	✓
Ep(0) uninfected cells	580,000	580,000	✓
Viral peak timing	Day 2 (± 6 h)	Days 1–4	✓
Peak viral load	8.9×10^6	10^5 – 10^8	✓
V at day 13.75	< 1% of peak	< 10% of peak	✓
Ep depletion at day 3	$\geq 9.3\%$	$\geq 5\%$	✓
n_{immune} at $t=0$	0 agents	0	✓
n_{immune} after day 1	≥ 1	> 0 within 1–4 d	✓
$n_{\text{immune}} \geq 0$ throughout	min = 0	≥ 0	✓
CC3D grid declared in ISSL	“90x90x2”	90x90x2	✓

bounds.

B. Validation Results

1) *Non-Invasive Adapter Verification*: All 67 automated tests pass. The Miao 2010 adapter initialises $T_{E_T} = 0$ and $k_E = 2 \times 10^{-5}$ (Table 1 constant), delegating all integration to roadrunner with no SBML edits. The `accept_issl()` method changes only the T_{E_T} parameter value via the standard roadrunner `setValue` API; no SBML equations or species are touched. The Sego 2020 CC3D adapter launch test confirms clean subprocess lifecycle management.

2) *Biological Plausibility of Individual Models*: ISSL-emitted quantities from each adapter were checked against published values (Table VI). All checks pass.

3) *Causal Ordering*: The coupled simulation generated ISSL checkpoint records across the full 14-day run (56 checkpoints) with no deadlocks, no temporal causality violations, and no watchdog alerts. The IPC handshake protocol — CC3D writes `abm_out.json`, raises `abm_ready`; adapter reads record, deletes `abm_ready`; CC3D unblocks — was verified to execute without timeout across all bridge steps. Checkpoint timestamps were strictly monotonically increasing. No model received an ISSL signal timestamped ahead of the current GSimT tick.

Figure 1(b) shows the GSimT timeline for a representative 5-day window. Causal-ordering annotation at the day 1 boundary confirms the `ABM.emit` $\rightarrow$ `ODE.accept` $\rightarrow$ `ODE.step` ordering. This ordering is enforced by the orchestrator’s topological DAG resolver and by the blocking IPC handshake, and was maintained without exception.

4) *Coupled Biological Dynamics*: The coupled simulation (Miao 2010 ODE + Sego 2020 CC3D ABM ensemble) ran for 14 days (56 GSimT checkpoints at 6 h intervals) with $N=5$ parallel CC3D instances. Runtime ensemble bounds are computed at each tick by the ensemble adapter and propagated through the ODE. Values below are median [ensemble range, $N=5$] from runtime ISSL records (Table VII).

Median peak $V = 9.03 \times 10^6$ copies/mL at day 2.25; median V at day 13.75 = 4.97×10^3 copies/mL = 0.06% of peak.

TABLE VII

14-DAY COUPLED TRAJECTORY. MEDIAN [ENSEMBLE RANGE, $N=5$]. DAY 0 CORRESPONDS TO THE STATE AFTER THE FIRST GSimT TICK (6 h OF ODE INTEGRATION FROM $V_0 = 1,000$ COPIES/ML).

GSimT	V median [range] (copies/mL)	n_{immune} [range]
Day 0	1.54×10^3	0 [0–0]
Day 1	3.29×10^5	6 [4–7]
Day 2	8.88×10^6	11 [10–12]
Day 3	6.27×10^6	15 [12–19]
Day 4	3.36×10^6	18 [14–24]
Day 5	1.78×10^6	23 [20–27]
Day 6	9.33×10^5	28 [26–32]
Day 7	4.85×10^5	30 [29–32]
Day 8	2.51×10^5	33 [32–39]
Day 9	1.29×10^5	39 [36–41]
Day 10	6.58×10^4	40 [37–43]
Day 11	3.34×10^4	44 [40–47]
Day 12	1.69×10^4	47 [44–51]
Day 13	8.43×10^3	50 [45–53]




Fig. 2. Coupled 14-day trajectory generated from ISSL JSON checkpoint files in `results/issl_14d/`. (a) Reference-replicate viral load V (log scale) and infected fraction $Eps/(Ep + Eps)$; viral peak annotated at day 2.25. (b) Median n_{immune} across $N=5$ replicates with IQR shading and individual replicate trajectories.

All 5 ensemble instances confirmed clearance ($V(\text{day } 14) < 0.1\%$ of peak). The ensemble bounds on n_{immune} are narrow on days 0–2 and widen through day 8, consistent with the Bernoulli seeding mechanism. The ODE viral load ensemble bounds (propagated via triple integration) track the n_{immune} uncertainty, demonstrating end-to-end UQ propagation.

Immune temporal lag. The n_{immune} trajectory shows first CC3D Immuncell agents appearing at day 1 (median = 6 agents), preceding the viral peak at day 2.25 by approximately 1.25 days. This rapid immune onset is consistent with the innate immune response timescale of 1–4 days post-infection [18] and emerges from the OISA-mediated signal routing.

CTL-mediated clearance. With n_{immune} growing from

TABLE VIII
SENSITIVITY OF COUPLED VIRAL TRAJECTORY TO κ AND SCALING
FACTOR. BOLD ROW: REFERENCE PARAMETER VALUES USED IN §V-B4.

κ	scaling	peak V	peak day	V_{14}/V_{peak}
3.5×10^{-8}	10	9.09×10^6	2.50	0.108%
3.5×10^{-8}	100	9.09×10^6	2.50	0.077%
3.5×10^{-8}	1000	9.09×10^6	2.50	0.003%
3.5×10^{-7}	10	9.08×10^6	2.50	0.105%
3.5×10^{-7}	100	9.01×10^6	2.50	0.055%
3.5×10^{-7}	1000	8.39×10^6	2.25	<0.001%
3.5×10^{-6}	10	9.08×10^6	2.50	0.102%
3.5×10^{-6}	100	8.94×10^6	2.50	0.044%
3.5×10^{-6}	1000	7.84×10^6	2.25	<0.001%

6 (day 1) to 50 (day 13), $T_{E_T} = n_{\text{immune}} \times 100$ grows from 600 to 5,000 CTL/mL in the Miao 2010 SBML, progressively amplifying the CTL killing term $k_E \cdot \text{Eps} \cdot T_{E_T}$. An isolated ODE run with $T_{E_T} = 0$ was confirmed to produce slower viral clearance, validating that the ABM $\rightarrow$ ODE signal pathway is a functionally significant contributor to the observed trajectory.

Real spatial ABM properties. The Immuncell agents are genuine spatial agents occupying sites on the $90 \times 90 \times 2$ Cellular Potts grid, subject to the Contact, Volume, and Chemotaxis plugins from the original Sego 2020 XML. Their recruitment is governed by ImmuneCellSeedingSteppable (stochastic) and movement by ChemotaxisSteppable (cytokine-gradient-directed). The n_{immune} values reported in the ISSL are counts from the live CC3D cell inventory (`cell.type == 5`), not from a closed-form approximation — the fundamental distinction from a scalar ODE surrogate.

5) *Parameter Sensitivity (ODE-side, coupling parameters only)*: This is an *approximate single-model* sensitivity sweep of the two coupling parameters only; the full ABM ensemble is not re-simulated for each grid cell. Sensitivity of the *model-internal* parameters of Miao 2010 (β_a , c_V , δ_{Es} , ...) and Sego 2020 (`ir_prob_scaling_factor`, `chemotaxis` coefficients, `max_ck_secrete_infect`, ...) is out of scope for this framework paper, since those parameters have already been calibrated against published experimental data by the original authors; a full coupled global sensitivity analysis (e.g. Sobol indices over both model-internal and coupling parameters) is deferred to future work.

The two empirically set coupling parameters — $\kappa = 3.5 \times 10^{-7}$ AU·mL/copies (ODE $\rightarrow$ ABM) and the scaling factor `_N_IMMUNE_TO_CTL_PER_ML` = 100 CTL/mL per CC3D agent (ABM $\rightarrow$ ODE) — were varied over ± 1 order of magnitude in a 3×3 grid (Table VIII). For each κ , the n_{immune} trajectory was *approximated* by shifting the reference ensemble median by the qualitatively observed ± 1 –2 day onset shift, rather than re-running the full $N = 5$ CC3D ensemble; the ODE was then integrated for 14 days with the shifted trajectory. An isolated ODE baseline ($T_{E_T} = 0$ throughout) is included for comparison (peak $V = 9.09 \times 10^6$, $V(\text{day } 14) = 1.02 \times 10^4$).

The sensitivity analysis reveals two findings. First, the scaling factor has the dominant effect: increasing it by $10 \times$

accelerates viral clearance by 2–3 orders of magnitude at day 14. This is expected because the scaling factor directly multiplies T_{E_T} in the CTL killing term. Second, κ has a secondary effect mediated through its influence on immune recruitment timing. Across all 9 combinations, peak V remains within $[7.84 \times 10^6, 9.09 \times 10^6]$ copies/mL — the same order of magnitude as the published Miao 2010 fit — and viral clearance ($V_{14} < 1\%$ of peak) is achieved in all cases. The qualitative trajectory shape is preserved across the full parameter grid, confirming that the coupled dynamics are robust to the coupling parameter choices within ± 1 order of magnitude.

VI. DISCUSSION

A. OISA and the CURE Guidelines

The CURE guidelines [5] are currently an arXiv preprint (peer-review status pending at time of submission); we interpret OISA’s design against their principles, supplemented by FAIR alignment where applicable. OISA maps cleanly to each CURE criterion: *credibility* via ISSL `ci_95` + ensemble bounds + the watchdog `divergence_score`; *provenance* via ISSL `envelope.model_version` (BIOMD0000000546 / commit-SHA + bridge tag for the ABM); *understandability* via `envelope` + `continuous_state` + `export_signals` per checkpoint; *reproducibility* via JSON-LD schema and commit-SHA `model_version`; *extensibility* via edge-wired config-graph substitutability; and *automated guideline checking* via constraint engine + watchdog at every GSimT tick. Three points deserve emphasis. First, OISA implements CURE credibility at the *composition* level: a composition can violate biological plausibility even when each constituent model is individually valid — e.g., a mis-specified tissue-to-systemic scaling factor silently miscalibrates CTL killing. The explicit `_N_IMMUNE_TO_CTL_PER_ML` = 100 parameter makes this decision auditable in the ISSL record. Second, `model_version` ties every checkpoint to a citable, immutable artefact. Third, zero internal biological modifications across two published models operationalises CURE extensibility at the multi-model scale.

B. Limitations

Validation scope. Validation is on a single organism (murine) and disease (influenza A); the coupled model has not been validated against in-vivo longitudinal tissue-level data, which are not published in directly comparable form.

Epithelial depletion. In Miao 2010, the high infection rate ($\beta_a = 10^{-6}$ mL·copies $^{-1}$ ·day $^{-1}$) drives near-complete epithelial depletion by day 2–3, after which clearance is driven by δ_{Es} and CTL killing. This is an intrinsic property of Miao 2010 and does not affect orchestration validation.

Feedback loops. The composition graph contains an ODE $\rightarrow$ ABM $\rightarrow$ ODE cycle. The 6 h back-edge delay is necessary for causal correctness and introduces a bounded lag that is negligible on influenza’s hour-to-day kinetic timescales.

ABM stochasticity and ensemble size. With $N = 5$, runtime bounds are sample min–max, not stable [2.5, 97.5] percentile estimates (which require $N \geq 20$). Production deployment requires a larger ensemble; wall-clock per tick is $\max(T_i)$ thanks to concurrent subprocesses.

Scaling factor and spatial scale mismatch. The `_N_IMMUNE_TO_CTL_PER_ML = 100` conversion is an approximation; more fundamentally, the Miao 2010 ODE ($V_{\text{total}} \approx 1 \text{ mL}$) and the Sego 2020 tissue patch ($\approx 10^{-3} \text{ mL}$) differ by ~ 3 OOM in volume. We omit volume normalisation because (i) the present target is orchestration correctness, not biological prediction, and (ii) the empirically tuned κ absorbs the mismatch implicitly. Predictive deployment must introduce explicit $V_{\text{tissue}} = V_{\text{systemic}} \times (V_{\text{patch}}/V_{\text{total}})$ normalisation and recalibrate κ and the scaling factor against paired tissue + systemic data.

IPC performance. The file-based IPC adds $\approx 50 \text{ ms}$ per tick ($< 1 \text{ s}$ total over 14 days), negligible vs. CC3D time. Sub-6-hour inter-model communication would require socket or shared-memory transport (2–3 OOM latency reduction).

VII. CONCLUSION

OISA delivers three contributions beyond the state of the art in multi-formalism immune model composition: (1) runtime UQ propagated at every checkpoint from a stochastic ABM ensemble through a deterministic ODE; (2) model-derived transfer lags on edges, demonstrated by a blood transit transfer model that computes signal delay dynamically; and (3) a biological plausibility constraint engine at the composition level. The architecture is validated by coupling Miao 2010 (SBML BIOMD0000000546) and the full Sego 2020 CC3D ABM (12 steppables, $90 \times 90 \times 2$ grid, @5b7e42c) with zero changes to either model. Over 14 days ($N = 5$ rolling ensemble, 56 GSimT checkpoints), ISSL records reproduce Miao 2010 viral kinetics (peak $V = 9.0 \times 10^6$ at day 2.25; $< 0.1\%$ of peak by day 13.75) and spatial immune recruitment (median $n_{\text{immune}} = 6$ [4–7] at day 1; 50 [40–53] at day 13). Immune temporal lag and CTL-mediated clearance acceleration emerge as bidirectional-coupling properties that neither model alone produces. Formalism-agnostic composition — including full spatial ABMs as independent processes with runtime UQ — is achievable through adapter-only integration with no model rewrites.

DATA, CODE, AND COMPETING INTERESTS

OISA specifications, ISSL JSON Schema (schemas/issl_v1.schema.json, Draft 2020-12), reference orchestrator, adapter code (~ 300 lines), configuration graph, 14-day run data (280 ISSL files, results/issl_14d/), sensitivity grid, figure scripts, and the full test suite (67 adapter tests + 51 paper-data consistency tests) are available at [repository URL to be added prior to submission]. Miao 2010 SBML is from BioModels unmodified; Sego 2020 is from covid-tissue-models@5b7e42c under CompuCell3D 4.8.0. The author declares no competing interests.

REFERENCES

- [1] M. Hucka, F. T. Bergmann, A. Dräger, S. Hoops, S. M. Keating, N. Le Novère, C. J. Myers, B. G. Olivier, S. Sahle, J. C. Schaff, L. P. Smith, D. Waltemath, and D. J. Wilkinson, “The systems biology markup language (SBML): Language specification for level 3 version 1 core,” *Journal of Integrative Bioinformatics*, vol. 12, no. 266, 2015.
- [2] M. Clerx, M. T. Cooling, J. Cooper, A. Garry, K. Moyle, D. P. Nickerson, P. M. F. Nielsen, and H. Sorby, “CellML 2.0,” *Journal of Integrative Bioinformatics*, vol. 17, no. 2–3, p. 20200021, 2020.
- [3] D. Waltemath, R. Adams, F. T. Bergmann, M. Hucka, F. Kolpakov, A. K. Miller, I. I. Moraru, D. Nickerson, S. Sahle, J. L. Snoep, and N. Le Novère, “Reproducible computational biology experiments with SED-ML,” *BMC Systems Biology*, vol. 5, p. 198, 2011.
- [4] E. Agmon, R. K. Spangler, C. J. Skalnik, W. Poole, S. M. Peirce, J. H. Morrison, and M. W. Covert, “Vivarium: an interface and engine for integrative multiscale modeling in computational biology,” *Bioinformatics*, vol. 38, pp. 1972–1979, 2022.
- [5] H. M. Sauro *et al.*, “From FAIR to CURE: Guidelines for computational models of biological systems,” arXiv:2502.15597, 2025, preprint; peer-review status pending at time of submission.
- [6] A. Ghaffarizadeh, R. Heiland, S. H. Friedman, S. M. Mumenthaler, and P. Macklin, “PhysiCell: An open source physics-based cell simulator for 3-D multicellular systems,” *PLoS Computational Biology*, vol. 14, p. e1005991, 2018.
- [7] M. Ponce-de Leon, A. Montagud, C. Akasiadis, J. Schreiber, T. Ntiniakou, and A. Valencia, “PhysiBoSS 2.0: a sustainable integration of stochastic Boolean and agent-based modelling frameworks,” *npj Systems Biology and Applications*, vol. 9, p. 54, 2023.
- [8] M. Getz, Y. Wang, G. An, A. Becker, C. Cockrell, N. Collier, M. Craig, C. L. Davis, J. Faeder, A. N. F. Versypt, J. F. Gianlupi, J. A. Glazier, S. Hamis, R. Heiland, T. Hillen, D. Hou, M. A. Nasr, P. A. Knapp, P. Macklin, T. Nguyen, M. Saglam, J. E. Shoemaker, A. M. Smith, J. J. F. Weaver, and P. Macklin, “Rapid community-driven development of a SARS-CoV-2 tissue simulator,” *iScience*, vol. 23, p. 101734, 2020.
- [9] T. J. Sego, J. O. Aponte-Serrano, J. F. Gianlupi, S. R. Heaps, K. Breithaupt, L. Brusch, J. M. Osborne, E. M. Quardokus, R. K. Plemper, and J. A. Glazier, “A modular framework for multiscale, multicellular, spatiotemporal modeling of acute primary viral infection and immune response in epithelial tissues and its application to drug therapy timing and effectiveness,” *PLoS Computational Biology*, vol. 16, no. 11, p. e1008451, 2020, GitHub: covid-tissue-models/covid-tissue-response-models @5b7e42c.
- [10] H. Miao, J. A. Hollenbaugh, M. S. Zand, J. Holden-Wiltse, T. R. Mosmann, A. S. Perelson, H. Wu, and D. J. Topham, “Quantifying the early immune response and adaptive immune response kinetics in mice infected with influenza A virus,” *Journal of Virology*, vol. 84, no. 14, pp. 7051–7062, 2010, bioModels: BIOMD0000000546.
- [11] F. T. Bergmann, R. Adams, S. Moodie, J. Cooper, M. Glont, M. Golebiewski, M. Hucka, C. Laibe, A. K. Miller, D. P. Nickerson, B. G. Olivier, N. Rodriguez, H. M. Sauro, M. Scharn, S. Soiland-Reyes, D. Waltemath, F. Yvon, and N. Le Novère, “COMBINE archive and OMEX format: one file to share all information to reproduce a modeling project,” *BMC Bioinformatics*, vol. 15, p. 369, 2014.
- [12] R. Laubenbacher, J. P. Sluka, and J. A. Glazier, “Building digital twins of the human immune system: toward a roadmap,” *npj Digital Medicine*, vol. 5, p. 64, 2022.
- [13] R. Laubenbacher, A. Niarakis, T. Helikar, G. An, B. Mehrad, C. Panayiotou, and J. A. Glazier, “Forum on immune digital twins: a meeting report,” *npj Systems Biology and Applications*, vol. 10, p. 19, 2024.
- [14] A. Niarakis, R. Laubenbacher, G. An, Y. Ilan, J. Fisher, Å. Flobak, K. Reiche, M. Rodríguez Martínez, L. Geris, L. Ladeira, L. Mortensen, B. M. Beckmann, F. Fröhlich, J. Honda, and T. Helikar, “Immune digital twins for complex human pathologies: applications, limitations, and challenges,” *npj Systems Biology and Applications*, vol. 10, p. 141, 2024.
- [15] Committee on Foundational Research Gaps and Future Directions for Digital Twins, *Foundational Research Gaps and Future Directions for Digital Twins*. Washington, DC: National Academies Press, 2023.
- [16] J. R. Karr, W. Liebermeister, A. P. Goldberg, J. A. P. Sekar, and B. Shaikh, “Model integration in computational biology: The role of reproducibility, credibility and utility,” *Frontiers in Systems Biology*, vol. 2, p. 822606, 2022.

- [17] M. Viceconti, M. De Vos, S. Mellone, and L. Geris, "From the digital twins in healthcare to the virtual human twin," *IEEE Journal of Biomedical and Health Informatics*, vol. 28, pp. 491–501, 2024.
- [18] A. Iwasaki and R. Medzhitov, "Regulation of adaptive immunity by the innate immune system," *Science*, vol. 327, no. 5963, pp. 291–295, 2010.